

Pthreads:

1.) Define thread ref. variables.

- Ex) pthread_t thread0;

2.) Create entry point function.

- function must have void* return type

- function must have a single void* arg.

- Ex) void* entry_func(void* param);

3.) Create thread.

- Create using pthread_create(thread, attr, start_routine, arg)

- thread: pointer to pthread_t variable

- attr: Extra attributes (can just set to NULL)

- start_routine: pointer to entry point function

- arg: pointer passed to entry point function as parameter.

- Ex) pthread_create(&thread0, NULL, entry_func, ¶meter);

* Might not need

4.) Join threads back together

- join using pthread_join(thread, status)

- thread: pthread_t variable used when thread created. (not a pointer this time)

- status: pointer to return value pointer (can just set to NULL)

- Ex) pthread_join(thread0, NULL);

Mutual Exclusion in Pthreads:

- Declare mutex variable: pthread_mutex_t mutex;

- Initialize mutex variable: pthread_mutex_init(&mutex, attr); (attr can be set to NULL for default)

- Lock + release when threads modify critical sections:

- Ex) pthread_mutex_lock(&mutex);

- critical_value += 1;

- pthread_mutex_unlock(&mutex);

MPI: Message Passing Interface;

~~Message Passing Interface~~

• Communicator & Group objects:

- define which collection of processes can communicate with each other.

- Use MPI_COMM_WORLD when communication required.

- Predefined communicator w/ all all MPI processes included.

• Rank ("task ID"):

- Unique integer identifier assigned when process initializes.

- Contiguous

- Begins at 0

- Used to specify source/destination of messages

• MPI Execution Environment Setup:

- MPI_Init(&argc, &argv):

• Initialize environment

• Must be called in every MPI program

• Must be called before any other MPI functions

• Must only be called once.

- MPI_Comm_rank(comm, &rank):

• Returns rank of a process in communicator

• @param rank = location to store returned rank number

- MPI_Comm_size(comm, &size):

• Returns the amount of processes that were requested for job

@param size = location to store returned number of processes

- MPI_Finalize():

• Terminates MPI execution environment

• Should be last MPI function called

• Point-to-Point Communication:

- Blocking Send: • MPI_Send(buffer, count, type, dest, tag, comm)

- Blocking Receive: • MPI_Recv(buffer, count, type, ~~dest~~ source, tag, comm, status)

MPI: Message Passing Interface(2):

• Point-to-Point Communication Cont'd:

- Message Passing Routine (MPI-Send/MPI-Recv) Parameters:

- Buffer: variable that is sent/received, usually passed by reference (&var)
- Count: number of data elements of "type" to be sent
- Type: MPI-CHAR, MPI-INT, MPI-LONG, MPI-FLOAT, etc (MPI-"type")

[MPI-Send] - Dest: Specifies the rank(Task ID) of the receiving process

[MPI-Recv] - Source: Specifies rank of sending process

- Tag: Any arbitrary non-negative number to identify message

- Comm: Communicator, usually MPI_COMM_WORLD

[MPI-Recv] - status: Indicates source of message and tag of message.

Consistency :

↳ Data-Centric Consistency Models :

↳ Sequential Consistency :

- When processes run concurrently, any valid interleaving of read/write operations is acceptable behavior, BUT all processes must see same interleaving of operations.
- Processes "see" writes from all other processes
- Process only "sees" its own reads
- Statements must be executed in program order

↳ Causal Consistency :

- "Weakened" form of sequential consistency.
 - Makes distinction between events that are potentially causally related and those that are not
- If event b is caused or influenced by event a, causality requires that everyone else first see a, then see b.
- Vector timestamps can be used to construct/maintain dependency graph

↳ Data-Centric Consistency w/ Grouping :

- Use shared synchronization variables to grant access to Critical Sections (CS).
- ENTER_CS → Acquire relevant synch. var.
- LEAVE_CS → Release synch. var.
- Possible for multiple processes to own some synch var but only in nonexclusive mode; can read, but not write, the associated data.
- Required Criteria:
 - At an acquire, all remote changes to guarded data must be made visible. (ie. released)
 - Before updating shared data, process must enter CS in exclusive mode; no other process may hold synch. var. in any (exclusive/nonexclusive) mode.
 - If process wants to enter CS in nonexclusive mode, it must fetch most recent copy of shared data from current owner of synch. var.

Client Centric Consistency Models:

↳ Eventual Consistency: * Is not a client-centric model

- In absence of frequent updates, all replicas of data set converge towards identical copies.
- Only require that update eventually propagates to all replicas.
- Problem arises if client accesses different data replicas after updating another replica.

— Alleviated through Client-Centric Consistency Models.

- Provides guarantees for a Single client concerning consistency of accesses to a data store. No guarantees for concurrent accesses by different clients.

↳ Monotonic Read Consistency:

- Guarantees that if a process has seen a value x at time t , it will never see an older version of x at later time.
- Interested in operations carried out by a single Processor P .

↳ Monotonic Write Consistency:

- A write operation by a process on data item x is completed before any successive write operations on x by the same process


↳ Read-Your-Writes Consistency:

- Effect of write operation by a process on x will always be seen by a successive read operation on x by same process

↳ Write-Follow-Read Consistency:

- Write op. by a process on x following a previous read op on x by the same process is guaranteed to take place on some/more recent value of x that was read.

4. Processing Times: A = 2, B = 4, C = 6
Every message takes + 1 (REQ, REL, GNT)



It is assumed that Process A&B perform calculations before a request. Each REQ/REL messages increments their clock by 1. PROC time is shown above for each process. Each GNT message sent from Process C has to increment the clock by 1. TOT is all the time taken from REQ to GNT.

Im not sure whether you need to process the last one. Unprocessed = 36, Processed = 42